

# Design tradeoffs in the Redis project

Salvatore Sanfilippo  
@antirez — Pivotal

# Before: some history...

- Software is about humans, not just engineering
- The history of an OSS project changes its design

# Too much time in my hands

- Creating a startup as a side project is hard.
- OSS is no exception, it is like a startup.
- You need a lot of time and focus.
- Good moments? After you graduate, or when your previous thing provided a *money buffer*.

# To be alone...

- It's hard to find a co-founder in an “pure” OSS project.
- To be alone is hard: accountability, self-motivation.
- To be alone is fun: single-minded design.
- To be alone narrows your POV: use the community.

# Design to scale yourself

- Complex software is rarely a good idea...
- Far worse if you want to maintain, often alone, a non trivial project.
- Sometimes complexity is intrinsically required: too simple is also not optimal.
- It is critical (and hard) to find a good balance.

# Important contribs?

## Mostly from payed people

- Peter Noordhuis (VMware)
- Matt Stancliff (Pivotal)
- Important fixes / debugging / contribs from RedisLabs, Weibo, Pinterest, Citrusbye.
- There are exceptions, of course (example: IPv6 support, high quality support in maling list).

# Stay motivated

- Rule 0: Do things you want to see existing.
- Rule 1: Evolve by sub-projects that are *interesting to you*.
- Rule 2: Create an economically sustainable model.
- Rule 3: Build a successful OSS, or abandon and try again.

# Design tradeoffs

To design means to sacrifice something



# Copy on write based persistence

Sacrifice memory  
for predictability and design freedom

# Append-only disk access

Sacrifice incremental updates  
to avoid disk seeks and corruptions

# Single thread model

Sacrifice single process scalability  
for big wins in simplicity and reliability

(spoiler: you need to scale cross-process, anyway...)

# Asynchronous replication

Sacrifice safety

for two order of magnitude more performances

(but it's not a black-or-white tradeoff: we now have  
async ACKs from slaves)

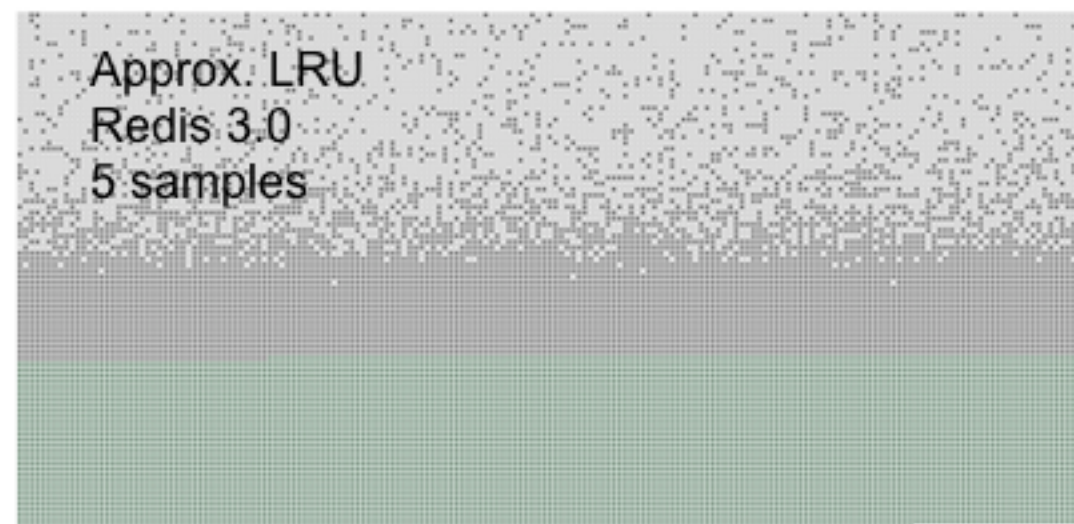

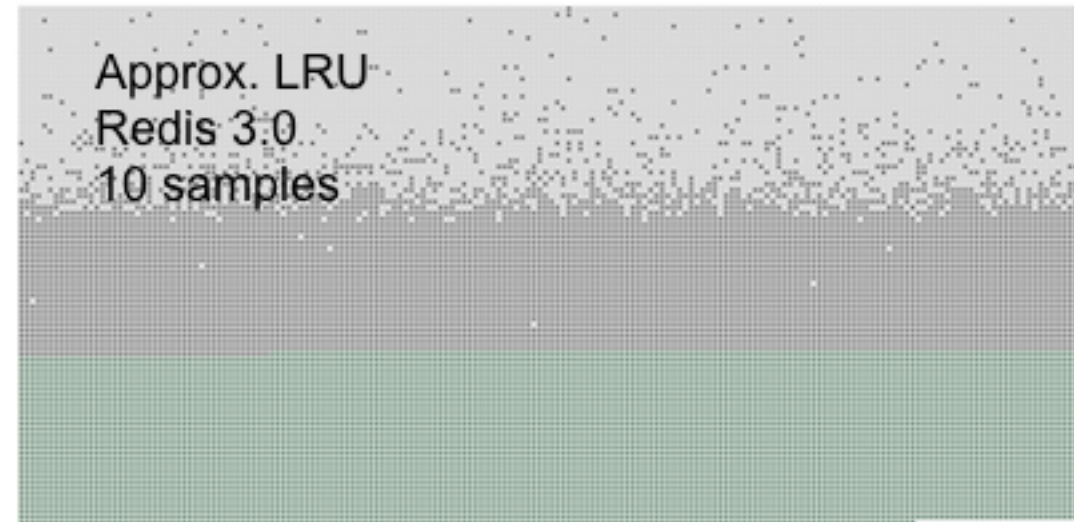

# No nested data structures

Sacrifice expressiveness  
for data model simplicity and simpler sharding

# Approximated LRU

Sacrifice theoretical correctness  
to save memory

We managed to improve it lately...



# No stored procedures

Sacrifice least surprise principle  
for a simpler to manage system



# Redis Cluster non trivial client role

Sacrifice separation of concerns  
for better latency and simplicity

# Redis Cluster consistency model

Sacrifice safety  
for data model freedom

(but limiting the danger when shit happens)

# Ask me anything!

Now or later...  
I'm @antirez on Twitter.